

# data-analytics

April 12, 2025

Q1.Data Cleaning and Imputation Techniques:

Load a dataset with missing values. Apply techniques like mean/mode/median imputation and compare the results.

```
[9]: !pip install pandas
```

```
Requirement already satisfied: pandas in c:\users\shinde
ankita\appdata\local\programs\python\python313\lib\site-packages (2.2.3)
Requirement already satisfied: numpy>=1.26.0 in c:\users\shinde
ankita\appdata\local\programs\python\python313\lib\site-packages (from pandas)
(2.2.4)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\shinde
ankita\appdata\local\programs\python\python313\lib\site-packages (from pandas)
(2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in c:\users\shinde
ankita\appdata\local\programs\python\python313\lib\site-packages (from pandas)
(2025.2)
Requirement already satisfied: tzdata>=2022.7 in c:\users\shinde
ankita\appdata\local\programs\python\python313\lib\site-packages (from pandas)
(2025.2)
Requirement already satisfied: six>=1.5 in c:\users\shinde
ankita\appdata\local\programs\python\python313\lib\site-packages (from python-
dateutil>=2.8.2->pandas) (1.17.0)
```

```
[6]: import pandas as pd
```

```
[7]: print(pd)
```

```
<module 'pandas' from 'C:\\Users\\Shinde
Ankita\\AppData\\Local\\Programs\\Python\\Python313\\Lib\\site-
packages\\pandas\\__init__.py'>
```

```
[12]: !pip install seaborn
```

```
Requirement already satisfied: seaborn in c:\users\shinde
ankita\appdata\local\programs\python\python313\lib\site-packages (0.13.2)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in c:\users\shinde
ankita\appdata\local\programs\python\python313\lib\site-packages (from seaborn)
```

(2.2.4)

Requirement already satisfied: pandas>=1.2 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from seaborn)

(2.2.3)

Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from seaborn)

(3.10.1)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from  
matplotlib!=3.6.1,>=3.4->seaborn) (1.3.1)

Requirement already satisfied: cycler>=0.10 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from  
matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from  
matplotlib!=3.6.1,>=3.4->seaborn) (4.57.0)

Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from  
matplotlib!=3.6.1,>=3.4->seaborn) (1.4.8)

Requirement already satisfied: packaging>=20.0 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from  
matplotlib!=3.6.1,>=3.4->seaborn) (24.2)

Requirement already satisfied: pillow>=8 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from  
matplotlib!=3.6.1,>=3.4->seaborn) (11.1.0)

Requirement already satisfied: pyparsing>=2.3.1 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from  
matplotlib!=3.6.1,>=3.4->seaborn) (3.2.3)

Requirement already satisfied: python-dateutil>=2.7 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from  
matplotlib!=3.6.1,>=3.4->seaborn) (2.9.0.post0)

Requirement already satisfied: pytz>=2020.1 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from  
pandas>=1.2->seaborn) (2025.2)

Requirement already satisfied: tzdata>=2022.7 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from  
pandas>=1.2->seaborn) (2025.2)

Requirement already satisfied: six>=1.5 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from python-  
dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.17.0)

```
[14]: import seaborn as sns
```

```
[15]: df = sns.load_dataset('titanic')
```

```
[16]: print(df.head())
```

```
survived  pclass      sex  age  sibsp  parch      fare embarked  class \
```

0	0	3	male	22.0	1	0	7.2500	S	Third
1	1	1	female	38.0	1	0	71.2833	C	First
2	1	3	female	26.0	0	0	7.9250	S	Third
3	1	1	female	35.0	1	0	53.1000	S	First
4	0	3	male	35.0	0	0	8.0500	S	Third

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True

```
[17]: print(df.isnull().sum())
```

```
survived      0
pclass        0
sex           0
age          177
sibsp         0
parch         0
fare          0
embarked      2
class         0
who           0
adult_male    0
deck         688
embark_town   2
alive         0
alone         0
dtype: int64
```

```
[18]: df_age = df[['age']].copy()
      print(df_age.describe())
```

```
count    age
count    714.000000
mean     29.699118
std      14.526497
min       0.420000
25%      20.125000
50%      28.000000
75%      38.000000
max       80.000000
```

(a) Mean Imputation

```
[20]: mean_value = df_age['age'].mean()
df_mean = df_age.fillna(mean_value)
```

(b) Median Imputation

```
[21]: median_value = df_age['age'].median()
df_median = df_age.fillna(median_value)
```

Mode Imputation

```
[22]: mode_value = df_age['age'].mode()[0]
df_mode = df_age.fillna(mode_value)
```

```
[23]: import matplotlib.pyplot as plt

plt.figure(figsize=(12, 6))

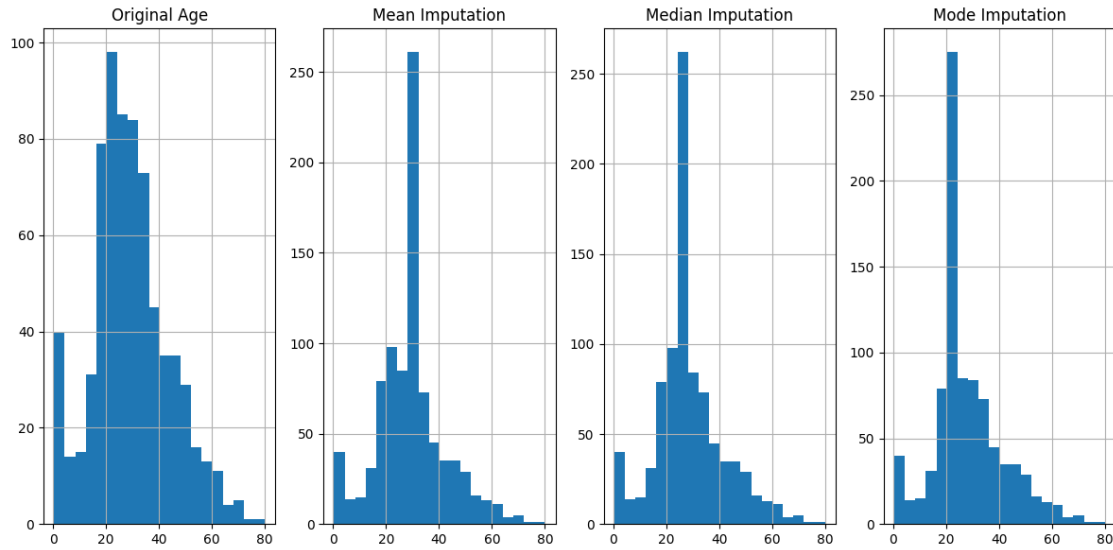
# Original (with missing values)
plt.subplot(1, 4, 1)
df_age['age'].hist(bins=20)
plt.title('Original Age')

# Mean Imputation
plt.subplot(1, 4, 2)
df_mean['age'].hist(bins=20)
plt.title('Mean Imputation')

# Median Imputation
plt.subplot(1, 4, 3)
df_median['age'].hist(bins=20)
plt.title('Median Imputation')

# Mode Imputation
plt.subplot(1, 4, 4)
df_mode['age'].hist(bins=20)
plt.title('Mode Imputation')

plt.tight_layout()
plt.show()
```



Q2.Data Analysis and Visualization: Use a dataset to: Plot scatterplots for numerical columns.Perform correlation analysis. Apply transformations (e.g., log, square root) and visualize the effect.

```
[24]: import seaborn as sns
import pandas as pd

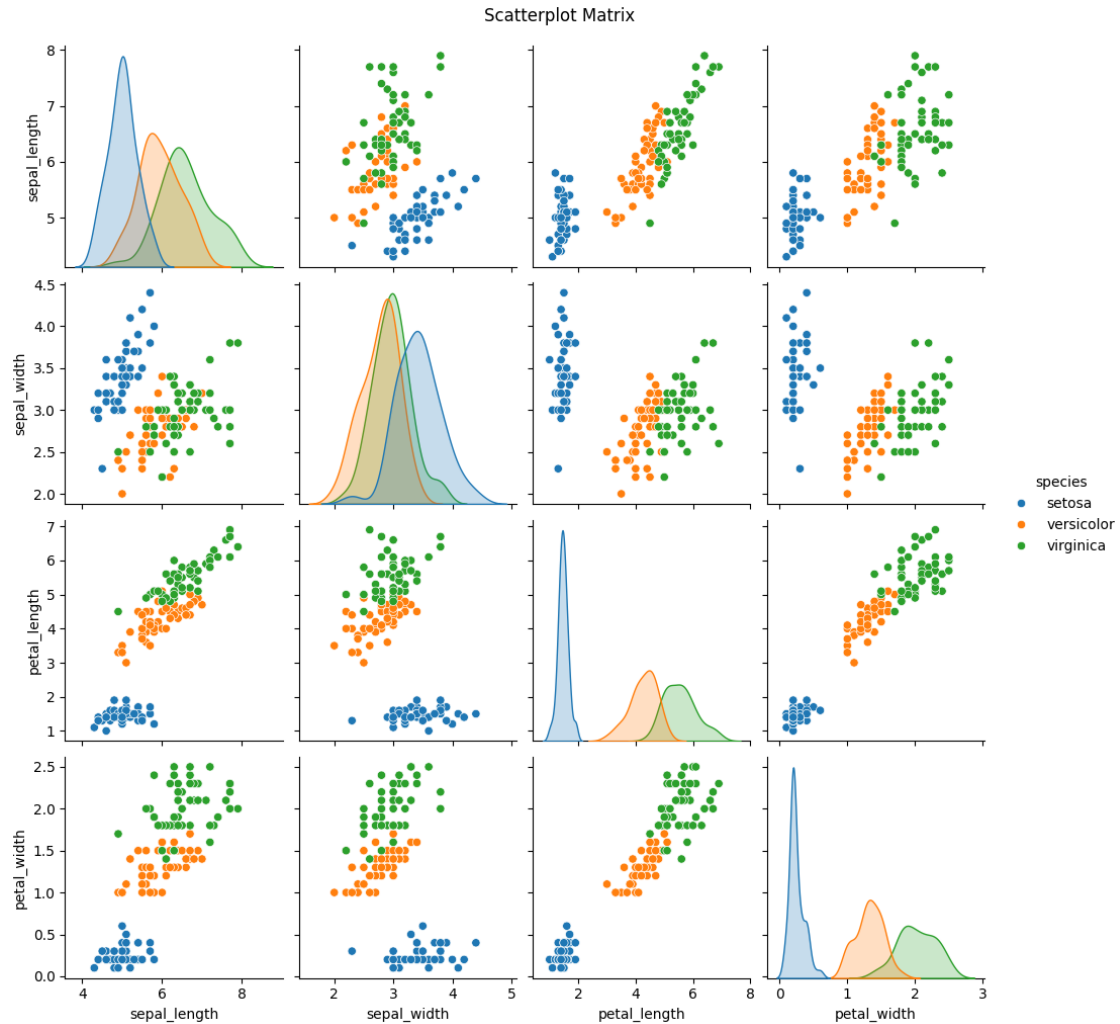
# Load Iris dataset
df = sns.load_dataset('iris')

# Show first few rows
print(df.head())
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa


```
[25]: import matplotlib.pyplot as plt

# Pairplot to plot all scatterplots between numerical columns
sns.pairplot(df, hue='species')
plt.suptitle("Scatterplot Matrix", y=1.02)
plt.show()
```



```
[26]: # Compute correlation
correlation_matrix = df.corr(numeric_only=True)

# Heatmap
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title("Correlation Heatmap")
plt.show()
```



```
[27]: import numpy as np


df['log_petal_length'] = np.log(df['petal_length'])

# Scatterplot before and after transformation
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
sns.scatterplot(x='petal_length', y='sepal_length', data=df)
plt.title("Original Petal Length")

plt.subplot(1, 2, 2)
sns.scatterplot(x='log_petal_length', y='sepal_length', data=df)
plt.title("Log-Transformed Petal Length")

plt.tight_layout()
plt.show()
```




```
[28]: df['sqrt_petal_length'] = np.sqrt(df['petal_length'])

plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
sns.scatterplot(x='sqrt_petal_length', y='sepal_length', data=df)
plt.title("Square Root Transformed Petal Length")

plt.tight_layout()
plt.show()
```





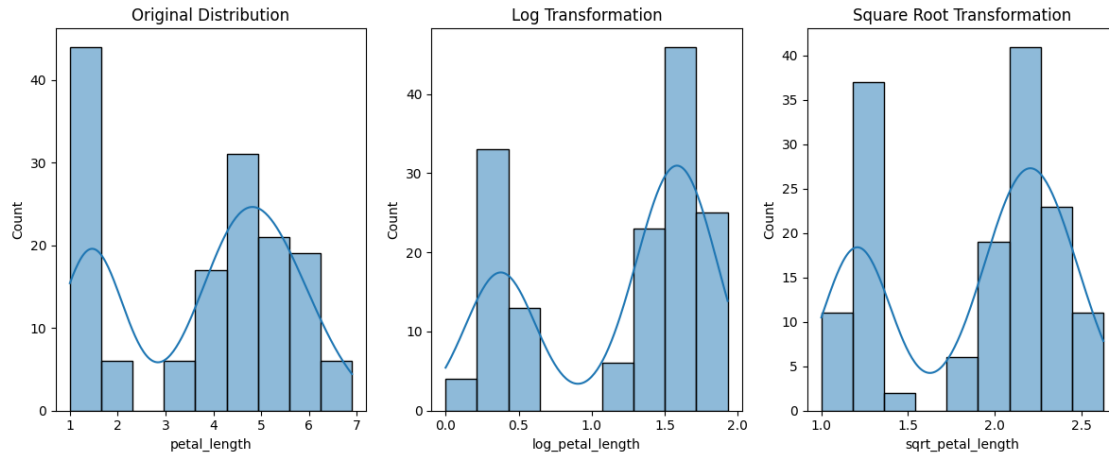
```
[29]: plt.figure(figsize=(12, 5))

# Original
plt.subplot(1, 3, 1)
sns.histplot(df['petal_length'], kde=True)
plt.title("Original Distribution")

# Log Transformed
plt.subplot(1, 3, 2)
sns.histplot(df['log_petal_length'], kde=True)
plt.title("Log Transformation")

# Sqrt Transformed
plt.subplot(1, 3, 3)
sns.histplot(df['sqrt_petal_length'], kde=True)
plt.title("Square Root Transformation")

plt.tight_layout()
plt.show()
```



### Q3.Encoding Methods:

Encode a categorical dataset using One-Hot Encoding and Label Encoding. Compare the effect of both methods on a machine-learning model.

Step 1: Load and Preprocess Titanic Dataset

```
[30]: import seaborn as sns
import pandas as pd

# Load dataset
df = sns.load_dataset('titanic')

# Keep relevant columns
df = df[['survived', 'sex', 'embarked', 'pclass', 'age']]

# Drop rows with missing values
df.dropna(inplace=True)

# Show dataset
print(df.head())
```

	survived	sex	embarked	pclass	age
0	0	male	S	3	22.0
1	1	female	C	1	38.0
2	1	female	S	3	26.0
3	1	female	S	1	35.0
4	0	male	S	3	35.0

Step 2: Label Encoding

```
[34]: !pip install scikit-learn
```

Collecting scikit-learn

```

Downloading scikit_learn-1.6.1-cp313-cp313-win_amd64.whl.metadata (15 kB)
Requirement already satisfied: numpy>=1.19.5 in c:\users\shinde
ankita\appdata\local\programs\python\python313\lib\site-packages (from scikit-
learn) (2.2.4)
Collecting scipy>=1.6.0 (from scikit-learn)
  Downloading scipy-1.15.2-cp313-cp313-win_amd64.whl.metadata (60 kB)
Collecting joblib>=1.2.0 (from scikit-learn)
  Downloading joblib-1.4.2-py3-none-any.whl.metadata (5.4 kB)
Collecting threadpoolctl>=3.1.0 (from scikit-learn)
  Downloading threadpoolctl-3.6.0-py3-none-any.whl.metadata (13 kB)
Downloading scikit_learn-1.6.1-cp313-cp313-win_amd64.whl (11.1 MB)
----- 0.0/11.1 MB ? eta -:--:--
----- 0.3/11.1 MB ? eta -:--:--
-- ----- 0.8/11.1 MB 2.6 MB/s eta 0:00:04
---- ----- 1.3/11.1 MB 2.8 MB/s eta 0:00:04
----- 2.1/11.1 MB 2.8 MB/s eta 0:00:04
----- 2.9/11.1 MB 3.1 MB/s eta 0:00:03
----- 3.9/11.1 MB 3.5 MB/s eta 0:00:03
----- 4.7/11.1 MB 3.5 MB/s eta 0:00:02
----- 6.0/11.1 MB 3.9 MB/s eta 0:00:02
----- 7.3/11.1 MB 4.2 MB/s eta 0:00:01
----- 8.9/11.1 MB 4.5 MB/s eta 0:00:01
----- 10.7/11.1 MB 4.9 MB/s eta 0:00:01
----- 11.1/11.1 MB 4.9 MB/s eta 0:00:00
Downloading joblib-1.4.2-py3-none-any.whl (301 kB)
Downloading scipy-1.15.2-cp313-cp313-win_amd64.whl (41.0 MB)
----- 0.0/41.0 MB ? eta -:--:--
- ----- 1.3/41.0 MB 7.6 MB/s eta 0:00:06
-- ----- 2.6/41.0 MB 6.8 MB/s eta 0:00:06
---- ----- 4.2/41.0 MB 7.1 MB/s eta 0:00:06
----- 6.0/41.0 MB 7.4 MB/s eta 0:00:05
----- 7.3/41.0 MB 7.2 MB/s eta 0:00:05
----- 8.7/41.0 MB 7.0 MB/s eta 0:00:05
----- 10.5/41.0 MB 7.0 MB/s eta 0:00:05
----- 12.3/41.0 MB 7.1 MB/s eta 0:00:05
----- 13.9/41.0 MB 7.3 MB/s eta 0:00:04
----- 15.5/41.0 MB 7.4 MB/s eta 0:00:04
----- 17.0/41.0 MB 7.4 MB/s eta 0:00:04
----- 18.1/41.0 MB 7.2 MB/s eta 0:00:04
----- 19.4/41.0 MB 7.0 MB/s eta 0:00:04
----- 19.9/41.0 MB 6.8 MB/s eta 0:00:04
----- 21.8/41.0 MB 6.8 MB/s eta 0:00:03
----- 23.3/41.0 MB 6.9 MB/s eta 0:00:03
----- 25.2/41.0 MB 7.0 MB/s eta 0:00:03
----- 26.2/41.0 MB 6.9 MB/s eta 0:00:03
----- 27.5/41.0 MB 6.9 MB/s eta 0:00:02
----- 29.1/41.0 MB 6.9 MB/s eta 0:00:02
----- 30.7/41.0 MB 6.9 MB/s eta 0:00:02

```

```

----- 32.2/41.0 MB 6.9 MB/s eta 0:00:02
----- 33.3/41.0 MB 6.9 MB/s eta 0:00:02
----- 34.3/41.0 MB 6.8 MB/s eta 0:00:01
----- 35.4/41.0 MB 6.7 MB/s eta 0:00:01
----- 37.0/41.0 MB 6.7 MB/s eta 0:00:01
----- 38.5/41.0 MB 6.8 MB/s eta 0:00:01
----- 39.1/41.0 MB 6.7 MB/s eta 0:00:01
----- 39.6/41.0 MB 6.6 MB/s eta 0:00:01
----- 40.4/41.0 MB 6.5 MB/s eta 0:00:01
----- 40.9/41.0 MB 6.3 MB/s eta 0:00:01
----- 41.0/41.0 MB 6.3 MB/s eta 0:00:00

```

Downloading threadpoolctl-3.6.0-py3-none-any.whl (18 kB)

Installing collected packages: threadpoolctl, scipy, joblib, scikit-learn

Successfully installed joblib-1.4.2 scikit-learn-1.6.1 scipy-1.15.2

threadpoolctl-3.6.0

```
[35]: from sklearn.preprocessing import LabelEncoder
      from sklearn.model_selection import train_test_split
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score
```

```
[37]: from sklearn.preprocessing import LabelEncoder
      from sklearn.model_selection import train_test_split
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score

      # Copy dataset
      df_label = df.copy()

      # Apply Label Encoding to categorical columns
      le = LabelEncoder()
      df_label['sex'] = le.fit_transform(df_label['sex'])      # male = 1, female = 0
      df_label['embarked'] = le.fit_transform(df_label['embarked'])

      # Train-Test Split
      X_label = df_label.drop('survived', axis=1)
      y_label = df_label['survived']
      X_train, X_test, y_train, y_test = train_test_split(X_label, y_label,
      test_size=0.2, random_state=42)

      # Model
      model_label = RandomForestClassifier(random_state=42)
      model_label.fit(X_train, y_train)
      pred_label = model_label.predict(X_test)

      # Accuracy
```

```
acc_label = accuracy_score(y_test, pred_label)
print("Label Encoding Accuracy:", acc_label)
```

Label Encoding Accuracy: 0.7692307692307693

Step 3: One-Hot Encoding

## 1 Copy dataset

```
df_onehot = df.copy()
```

## 2 Apply One-Hot Encoding

```
df_onehot = pd.get_dummies(df_onehot, columns=['sex', 'embarked'], drop_first=True)
```

## 3 Train-Test Split

```
X_onehot = df_onehot.drop('survived', axis=1) y_onehot = df_onehot['survived'] X_train_oh,
X_test_oh, y_train_oh, y_test_oh = train_test_split(X_onehot, y_onehot, test_size=0.2, random_state=42)
```

## 4 Model

```
model_onehot = RandomForestClassifier(random_state=42) model_onehot.fit(X_train_oh,
y_train_oh) pred_onehot = model_onehot.predict(X_test_oh)
```

## 5 Accuracy

```
acc_onehot = accuracy_score(y_test_oh, pred_onehot) print("One-Hot Encoding Accuracy:",
acc_onehot)
```

Q4. Outlier Detection:

Use the Isolation Forest algorithm to detect and visualize outliers in a dataset.

Step 1: Import Libraries & Load Dataset

```
[40]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.ensemble import IsolationForest

# Load the Iris dataset
df = sns.load_dataset('iris')

# We'll only use numerical features
df = df[['sepal_length', 'sepal_width']]
```

```
# Display the data
print(df.head())
```

	sepal_length	sepal_width
0	5.1	3.5
1	4.9	3.0
2	4.7	3.2
3	4.6	3.1
4	5.0	3.6

Step 2: Apply Isolation Forest

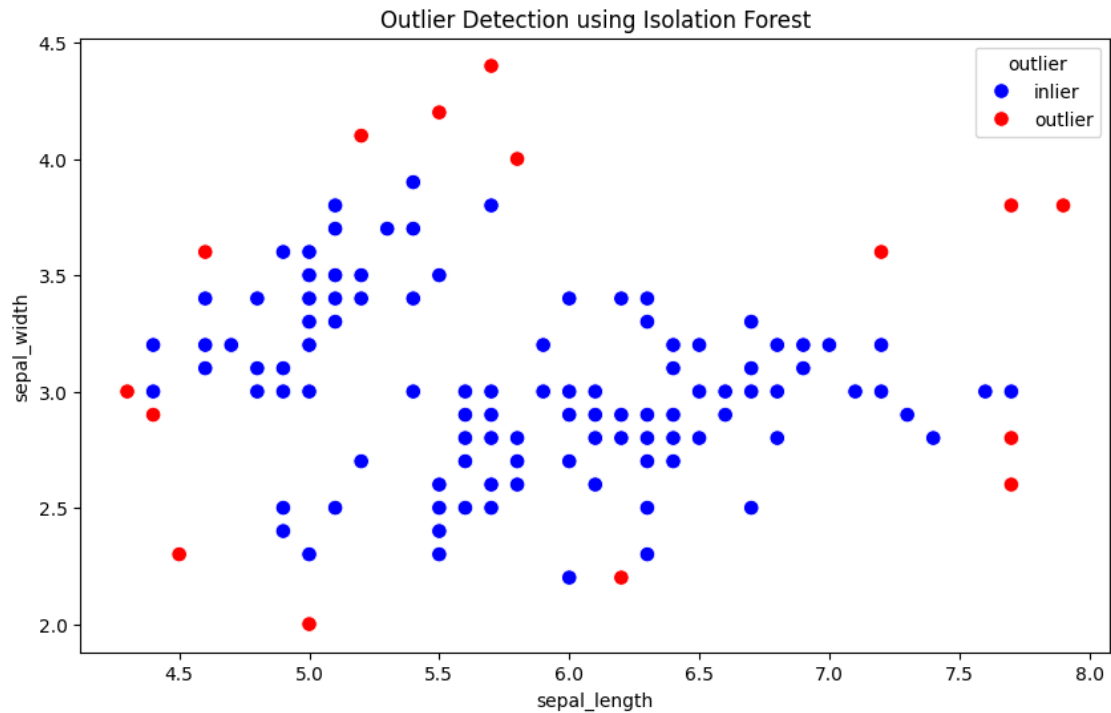
```
[41]: # Initialize the model
iso = IsolationForest(contamination=0.1, random_state=42) # 10% outliers

# Fit and predict
df['outlier'] = iso.fit_predict(df)

# Map results: -1 is outlier, 1 is inlier
df['outlier'] = df['outlier'].map({1: 'inlier', -1: 'outlier'})
```

Step 3: Visualize Outliers

```
[42]: # Scatter plot showing outliers
plt.figure(figsize=(10, 6))
sns.scatterplot(data=df, x='sepal_length', y='sepal_width', hue='outlier',
               palette={'inlier': 'blue', 'outlier': 'red'}, s=70)
plt.title("Outlier Detection using Isolation Forest")
plt.show()
```



Q5. Predictive Power Score (PPS):

Calculate the PPS for a dataset and interpret which variables are most predictive.

```
[ ]: !pip install ppscore
```

```
[ ]: import ppscore as pps
```

```
[ ]: import pandas as pd
import seaborn as sns
import ppscore as pps

# Load dataset
df = sns.load_dataset("titanic")

# Clean up: drop rows with NaNs and irrelevant columns
df = df[['survived', 'pclass', 'sex', 'age', 'sibsp', 'parch', 'fare', '
↳ 'embarked']].dropna()
```

```
[ ]: # PPS matrix: predictive power of each feature → target
pps_matrix = pps.matrix(df)

# Show top predictive features
pps_matrix_filtered = pps_matrix[pps_matrix['y'] == 'survived'].
↳ sort_values('ppscore', ascending=False)
```

```
print(pps_matrix_filtered[['x', 'y', 'ppscore']])
```

```
[ ]: # Top features predictive of 'survived'
top_predictors = pps_matrix[pps_matrix['y'] == 'survived'].
    ↪sort_values('ppscore', ascending=False)
print(top_predictors[['x', 'ppscore']])
```

```
[ ]: import pandas as pd

# Replace 'your_file.csv' with your actual file path
df = pd.read_csv("sample.csv")

# Preview it
df.head()
```

```
[ ]: !pip install ppscore
import ppscore as pps
```

```
[ ]: # Calculate PPS matrix
pps_matrix = pps.matrix(df)

# See top predictors for a specific column
target_column = 'your_target_column' # <- Replace this
pps_matrix[pps_matrix['y'] == target_column].sort_values('ppscore',
    ↪ascending=False)
```

```
[ ]: import seaborn as sns
import matplotlib.pyplot as plt

# Pivot for heatmap
pps_heatmap = pps_matrix.pivot(index='x', columns='y', values='ppscore')

plt.figure(figsize=(12, 6))
sns.heatmap(pps_heatmap, annot=True, cmap="YlGnBu")
plt.title("Predictive Power Score Matrix")
plt.show()
```

#### Q6.Simple and Multiple Linear Regression:

Implement simple and multiple linear regression on a dataset. Evaluate the model's performance using R- squared and mean squared error (MSE).

Step 1: Import Libraries & Dataset

```
[4]: import pandas as pd
import numpy as np
import seaborn as sns
```



```

import matplotlib.pyplot as plt

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score, mean_squared_error

# Load dataset (we'll use 'mpg' from seaborn as it's perfect)
df = sns.load_dataset('mpg').dropna()

# Preview
df.head()

```

```

[4]:      mpg  cylinders  displacement  horsepower  weight  acceleration  \
0   18.0          8         307.0         130.0    3504          12.0
1   15.0          8         350.0         165.0    3693          11.5
2   18.0          8         318.0         150.0    3436          11.0
3   16.0          8         304.0         150.0    3433          12.0
4   17.0          8         302.0         140.0    3449          10.5

      model_year origin          name
0           70    usa  chevrolet chevelle malibu
1           70    usa      buick skylark 320
2           70    usa    plymouth satellite
3           70    usa      amc rebel sst
4           70    usa      ford torino

```

Step 2: Simple Linear Regression

```

[5]: # Simple Linear Regression
X_simple = df[['horsepower']]
y = df['mpg']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X_simple, y, test_size=0.2,
                                                    random_state=42)

# Model
model_simple = LinearRegression()
model_simple.fit(X_train, y_train)

# Predict
y_pred = model_simple.predict(X_test)

# Evaluation
r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)

```

```
print("Simple Linear Regression")
print("R-squared:", round(r2, 3))
print("Mean Squared Error:", round(mse, 3))
```

Simple Linear Regression  
R-squared: 0.566  
Mean Squared Error: 22.153

Step 3: Multiple Linear Regression

```
[6]: # Multiple Linear Regression
features = ['horsepower', 'weight', 'acceleration', 'displacement']
X_multi = df[features]

# Split
X_train, X_test, y_train, y_test = train_test_split(X_multi, y, test_size=0.2,
                                                    random_state=42)

# Model
model_multi = LinearRegression()
model_multi.fit(X_train, y_train)

# Predict
y_pred_multi = model_multi.predict(X_test)

# Evaluation
r2_multi = r2_score(y_test, y_pred_multi)
mse_multi = mean_squared_error(y_test, y_pred_multi)

print("\nMultiple Linear Regression")
print("R-squared:", round(r2_multi, 3))
print("Mean Squared Error:", round(mse_multi, 3))
```

Multiple Linear Regression  
R-squared: 0.646  
Mean Squared Error: 18.066

Q7. Build a logistic regression model to classify binary outcomes (e.g., predicting if a customer will buy a product). Evaluate the model using confusion matrix metrics.

Step 1: Import Libraries & Load Dataset

```
[7]: import pandas as pd
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
```

```
# Load Titanic dataset
df = sns.load_dataset('titanic')

# Select useful features and drop missing data
df = df[['survived', 'pclass', 'sex', 'age', 'fare']].dropna()
df['sex'] = df['sex'].map({'male': 0, 'female': 1}) # Encode 'sex'
```

Step 2: Split Data

```
[8]: # Features and target
X = df[['pclass', 'sex', 'age', 'fare']]
y = df['survived']

# Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

Step 3: Train Logistic Regression

```
[10]: model = LogisticRegression()
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)
```

Step 4: Evaluate Using Confusion Matrix

```
[11]: from sklearn.metrics import ConfusionMatrixDisplay


# Confusion Matrix & Classification Report
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))

# Optional visualization
ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
```

```
[[68 19]
 [16 40]]
```

	precision	recall	f1-score	support
0	0.81	0.78	0.80	87
1	0.68	0.71	0.70	56
accuracy			0.76	143
macro avg	0.74	0.75	0.75	143
weighted avg	0.76	0.76	0.76	143

```
[11]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1d2eb886900>
```



Q8.Clustering Techniques: Perform K-Means and hierarchical clustering on a dataset. Visualize the clusters and interpret the results.

Step 1: Import Libraries & Load Data

```
[12]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from scipy.cluster.hierarchy import linkage, dendrogram, fcluster
from sklearn.preprocessing import StandardScaler

# Load Iris dataset
df = sns.load_dataset("iris")
df.head()
```

```
[12]:   sepal_length  sepal_width  petal_length  petal_width  species
0          5.1          3.5          1.4          0.2    setosa
1          4.9          3.0          1.4          0.2    setosa
2          4.7          3.2          1.3          0.2    setosa
```

3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

Step 2: Prepare Features (Drop the label)

```
[13]: X = df.drop('species', axis=1)


# Standardize features (important for distance-based algorithms)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

Step 3: Apply K-Means

```
[14]: kmeans = KMeans(n_clusters=3, random_state=42)
df['kmeans_cluster'] = kmeans.fit_predict(X_scaled)
```

Step 4: Visualize K-Means Clusters

```
[15]: sns.scatterplot(data=df, x='sepal_length', y='sepal_width',
    ↪ hue='kmeans_cluster', palette='viridis')
plt.title("K-Means Clustering")
plt.show()
```




## Hierarchical Clustering Step 5: Apply Hierarchical Clustering

```
[16]: linkage_matrix = linkage(X_scaled, method='ward')
```

## Step 6: Dendrogram

```
[18]: plt.figure(figsize=(10, 5))
dendrogram(linkage_matrix, truncate_mode='lastp', p=30)
plt.title("Hierarchical Clustering Dendrogram")
plt.xlabel("Sample index")
plt.ylabel("Distance")
plt.show()
```




## Step 7: Assign Cluster Labels

```
[19]: df['hierarchical_cluster'] = fcluster(linkage_matrix, t=3, criterion='maxclust')
```

## Step 8: Visualize Hierarchical Clusters

```
[20]: sns.scatterplot(data=df, x='sepal_length', y='sepal_width',
hue='hierarchical_cluster', palette='coolwarm')
plt.title("Hierarchical Clustering")
plt.show()
```



Q9.Principal Component Analysis (PCA):

Apply PCA on a high-dimensional dataset. Reduce the dimensions and visualize the transformed data.

Step 1: Import Libraries & Load Data

```
[21]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

# Load dataset
df = sns.load_dataset('iris')

# Separate features and label
X = df.drop('species', axis=1)
y = df['species']
```

Step 2: Standardize the Data

```
[22]: scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```


Step 3: Apply PCA

```
[23]: pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

# Convert to DataFrame
df_pca = pd.DataFrame(X_pca, columns=['PC1', 'PC2'])
df_pca['species'] = y
```

Step 4: Visualize PCA-Reduced Data

```
[25]: sns.scatterplot(data=df_pca, x='PC1', y='PC2', hue='species', palette='Set2')
plt.title("PCA: Iris Dataset (2D Projection)")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.grid(True)
plt.show()
```





Unit 3: Specialized Applications in AI and ML Market Basket Analysis: Q10.Implement Association Rule Mining using the Apriori algorithm. Identify frequent itemsets and generate association rules for a transactional dataset.

[26]: `!pip install mlxtend`

Collecting mlxtend

Downloading mlxtend-0.23.4-py3-none-any.whl.metadata (7.3 kB)

Requirement already satisfied: scipy>=1.2.1 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from mlxtend) (1.15.2)

Requirement already satisfied: numpy>=1.16.2 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from mlxtend) (2.2.4)

Requirement already satisfied: pandas>=0.24.2 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from mlxtend) (2.2.3)

Requirement already satisfied: scikit-learn>=1.3.1 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from mlxtend) (1.6.1)

Requirement already satisfied: matplotlib>=3.0.0 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from mlxtend) (3.10.1)

Requirement already satisfied: joblib>=0.13.2 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from mlxtend) (1.4.2)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (1.3.1)

Requirement already satisfied: cyclor>=0.10 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (4.57.0)

Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (1.4.8)

Requirement already satisfied: packaging>=20.0 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (24.2)

Requirement already satisfied: pillow>=8 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (11.1.0)

Requirement already satisfied: pyparsing>=2.3.1 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (3.2.3)

Requirement already satisfied: python-dateutil>=2.7 in c:\users\shinde

```

ankita\appdata\local\programs\python\python313\lib\site-packages (from
matplotlib>=3.0.0->mlxtend) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in c:\users\shinde
ankita\appdata\local\programs\python\python313\lib\site-packages (from
pandas>=0.24.2->mlxtend) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in c:\users\shinde
ankita\appdata\local\programs\python\python313\lib\site-packages (from
pandas>=0.24.2->mlxtend) (2025.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in c:\users\shinde
ankita\appdata\local\programs\python\python313\lib\site-packages (from scikit-
learn>=1.3.1->mlxtend) (3.6.0)
Requirement already satisfied: six>=1.5 in c:\users\shinde
ankita\appdata\local\programs\python\python313\lib\site-packages (from python-
dateutil>=2.7->matplotlib>=3.0.0->mlxtend) (1.17.0)
Downloading mlxtend-0.23.4-py3-none-any.whl (1.4 MB)
----- 0.0/1.4 MB ? eta -:-:--
----- 0.8/1.4 MB 8.0 MB/s eta 0:00:01
----- 1.4/1.4 MB 6.7 MB/s eta 0:00:00

Installing collected packages: mlxtend
Successfully installed mlxtend-0.23.4

```

```

[27]: import pandas as pd
      from mlxtend.preprocessing import TransactionEncoder

      # Sample transactional data
      transactions = [
          ['milk', 'bread', 'butter'],
          ['bread', 'diapers', 'beer', 'eggs'],
          ['milk', 'diapers', 'beer', 'cola'],
          ['bread', 'milk', 'diapers', 'beer'],
          ['bread', 'milk', 'diapers', 'cola']
      ]

      # Convert to one-hot encoding format
      te = TransactionEncoder()
      te_ary = te.fit(transactions).transform(transactions)
      df = pd.DataFrame(te_ary, columns=te.columns_)
      df.head()

```

```

[27]:
   beer  bread  butter  cola  diapers  eggs  milk
0  False   True   True  False   False  False  True
1   True   True  False  False    True   True  False
2   True  False  False   True    True  False  True
3   True   True  False  False    True  False  True
4  False   True  False   True    True  False  True

```

```
[28]: from mlxtend.frequent_patterns import apriori
```

```
# Get frequent itemsets
frequent_itemsets = apriori(df, min_support=0.6, use_colnames=True)
frequent_itemsets
```

```
[28]:
```

	support	itemsets
0	0.6	(beer)
1	0.8	(bread)
2	0.8	(diapers)
3	0.8	(milk)
4	0.6	(diapers, beer)
5	0.6	(bread, diapers)
6	0.6	(milk, bread)
7	0.6	(milk, diapers)

```
[29]: from mlxtend.frequent_patterns import association_rules
```

```
# Generate rules
rules = association_rules(frequent_itemsets, metric="lift", min_threshold=1.0)
rules[['antecedents', 'consequents', 'support', 'confidence', 'lift']]
```

```
[29]:
```

	antecedents	consequents	support	confidence	lift
0	(diapers)	(beer)	0.6	0.75	1.25
1	(beer)	(diapers)	0.6	1.00	1.25

#### Q11.Recommendation Systems:

Build a collaborative filtering recommendation system using a movie or product dataset. Compare results using user-based and item-based filtering.

```
[31]: import pandas as pd
from sklearn.metrics.pairwise import cosine_similarity

# Sample ratings data (user-movie matrix)
ratings_dict = {
    'User': ['A', 'A', 'A', 'B', 'B', 'C', 'C', 'D', 'E'],
    'Movie': ['Titanic', 'Avatar', 'Avengers', 'Titanic', 'Avengers', 'Avatar',
    ↪ 'Avengers', 'Titanic', 'Avatar'],
    'Rating': [5, 3, 4, 4, 5, 2, 4, 5, 3]
}

df = pd.DataFrame(ratings_dict)
ratings_matrix = df.pivot_table(index='User', columns='Movie', values='Rating').
    ↪ fillna(0)
ratings_matrix
```

```
[31]: Movie  Avatar  Avengers  Titanic
      User
A         3.0      4.0      5.0
B         0.0      5.0      4.0
C         2.0      4.0      0.0
D         0.0      0.0      5.0
E         3.0      0.0      0.0
```

```
[32]: # Compute cosine similarity between users
      user_similarity = cosine_similarity(ratings_matrix)
      user_sim_df = pd.DataFrame(user_similarity, index=ratings_matrix.index,
      ↪columns=ratings_matrix.index)
      user_sim_df
```

```
[32]: User      A      B      C      D      E
      User
A    1.000000  0.883452  0.695701  0.707107  0.424264
B    0.883452  1.000000  0.698430  0.624695  0.000000
C    0.695701  0.698430  1.000000  0.000000  0.447214
D    0.707107  0.624695  0.000000  1.000000  0.000000
E    0.424264  0.000000  0.447214  0.000000  1.000000
```

```
[33]: # Transpose and compute cosine similarity between movies
      item_similarity = cosine_similarity(ratings_matrix.T)
      item_sim_df = pd.DataFrame(item_similarity, index=ratings_matrix.columns,
      ↪columns=ratings_matrix.columns)
      item_sim_df
```

```
[33]: Movie      Avatar  Avengers  Titanic
      Movie
Avatar    1.000000  0.564782  0.393648
Avengers  0.564782  1.000000  0.652155
Titanic   0.393648  0.652155  1.000000
```

```
[34]: similar_users = user_sim_df['C'].sort_values(ascending=False)[1:]
      top_user = similar_users.index[0]
      recommendation = ratings_matrix.loc[top_user][ratings_matrix.loc['C'] == 0].
      ↪sort_values(ascending=False)
      recommendation
```

```
[34]: Movie
      Titanic    4.0
      Name: B, dtype: float64
```

```
[35]: item_sim_df['Avatar'].sort_values(ascending=False)[1:]
```

```
[35]: Movie
      Avengers    0.564782
      Titanic    0.393648
      Name: Avatar, dtype: float64
```

## Q12.Tree-Based Feature Engineering

Apply tree-based methods to rank feature importance in a dataset. Use the results to train a simplified model.

```
[36]: import pandas as pd
import seaborn as sns
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt

# Load Iris dataset
df = sns.load_dataset("iris")

# Features and target
X = df.drop("species", axis=1)
y = df["species"]


# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)
```

```
[37]: # Train a random forest
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# Get feature importances
importances = rf.feature_importances_
feature_ranking = pd.Series(importances, index=X.columns).
    sort_values(ascending=False)
print(feature_ranking)
```

```
petal_length    0.439994
petal_width     0.421522
sepal_length    0.108098
sepal_width     0.030387
dtype: float64
```

```
[38]: feature_ranking.plot(kind='barh', title='Feature Importance (Random Forest)',
    color='skyblue')
plt.gca().invert_yaxis()
plt.show()
```



```
[39]: top_features = feature_ranking.head(2).index.tolist()

# Train with only top 2 features
X_train_simple = X_train[top_features]
X_test_simple = X_test[top_features]

rf_simple = RandomForestClassifier(n_estimators=100, random_state=42)
rf_simple.fit(X_train_simple, y_train)

# Evaluate
y_pred_simple = rf_simple.predict(X_test_simple)
accuracy = accuracy_score(y_test, y_pred_simple)
print("Simplified Model Accuracy:", round(accuracy * 100, 2), "%")
```

Simplified Model Accuracy: 100.0 %

#### Q13. Recursive Feature Elimination (RFE)

Perform feature selection using RFE. Evaluate the performance of a machine-learning model before and after feature selection.

```
[40]: import pandas as pd
import seaborn as sns
```

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.feature_selection import RFE

# Load dataset
df = sns.load_dataset('iris')

# Prepare features and target
X = df.drop('species', axis=1)
y = df['species']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

```

```

[41]: model = LogisticRegression(max_iter=200)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

# Accuracy before feature selection
base_acc = accuracy_score(y_test, y_pred)
print("Accuracy (All Features):", round(base_acc * 100, 2), "%")

```

Accuracy (All Features): 100.0 %

```

[42]: # Recursive Feature Elimination
selector = RFE(estimator=LogisticRegression(max_iter=200),
↳n_features_to_select=2)
selector.fit(X_train, y_train)

# Get selected features
selected_features = X.columns[selector.support_]
print("Selected Features by RFE:", selected_features.tolist())

```

Selected Features by RFE: ['petal\_length', 'petal\_width']

```

[43]: # Train with selected features only
X_train_rfe = X_train[selected_features]
X_test_rfe = X_test[selected_features]

model_rfe = LogisticRegression(max_iter=200)
model_rfe.fit(X_train_rfe, y_train)
y_pred_rfe = model_rfe.predict(X_test_rfe)

# Accuracy after RFE
rfe_acc = accuracy_score(y_test, y_pred_rfe)

```

```
print("Accuracy (After RFE):", round(rfe_acc * 100, 2), "%")
```

Accuracy (After RFE): 100.0 %

#### Q14.Train-Test Split and Cross-Validation

Split a dataset into train-test sets. Use Shuffle Cross- Validation to evaluate a model and compare the results.

```
[45]: import pandas as pd
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, ShuffleSplit, \
    cross_val_score
from sklearn.metrics import accuracy_score

# Load dataset
df = sns.load_dataset('iris')

# Features and target
X = df.drop('species', axis=1)
y = df['species']

[46]: # Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, \
    random_state=42)

# Train model
model = LogisticRegression(max_iter=200)
model.fit(X_train, y_train)

# Evaluate on test set
y_pred = model.predict(X_test)
acc_test = accuracy_score(y_test, y_pred)
print("Accuracy (Train-Test Split):", round(acc_test * 100, 2), "%")
```

Accuracy (Train-Test Split): 100.0 %

```
[47]: from sklearn.model_selection import ShuffleSplit

# Set up ShuffleSplit
shuffle_split = ShuffleSplit(n_splits=5, test_size=0.2, random_state=42)

# Cross-validation scores
cv_scores = cross_val_score(model, X, y, cv=shuffle_split)

print("Cross-Validation Scores:", cv_scores)
print("Mean CV Accuracy:", round(cv_scores.mean() * 100, 2), "%")
```



Cross-Validation Scores: [1. 0.96666667 0.96666667 0.93333333  
0.93333333]  
Mean CV Accuracy: 96.0 %

#### Q15. Bagging and Random Forest

Build and evaluate a Random Forest model. Visualize the decision trees and feature importance.

```
[48]: import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

# Load the Iris dataset
df = sns.load_dataset('iris')
X = df.drop('species', axis=1)
y = df['species']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)
```

```
[49]: # Train a Random Forest classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)


# Predict and evaluate
y_pred = rf.predict(X_test)
acc = accuracy_score(y_test, y_pred)
print("Random Forest Accuracy:", round(acc * 100, 2), "%")
```

Random Forest Accuracy: 100.0 %

```
[50]: import matplotlib.pyplot as plt

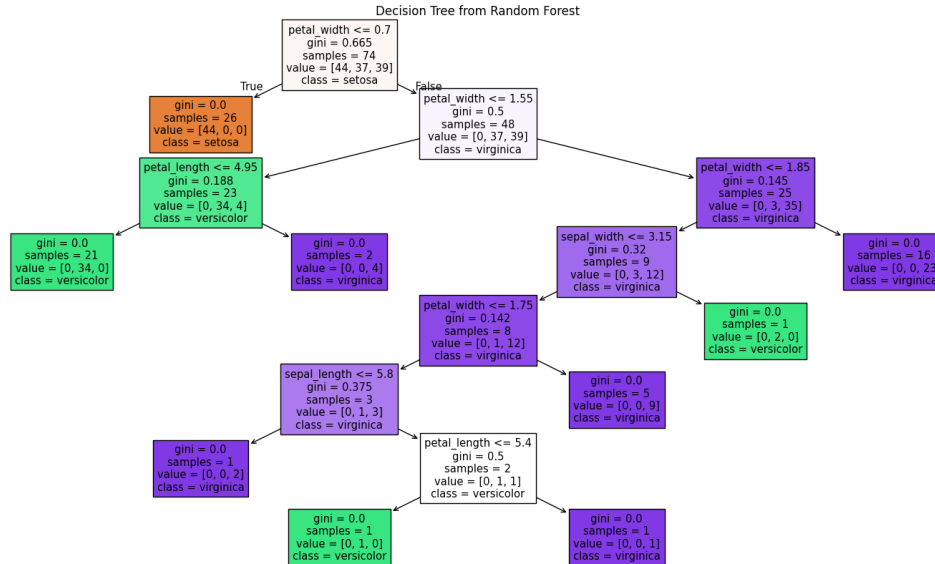
# Get and plot feature importances
importances = rf.feature_importances_
feat_names = X.columns
feat_imp_df = pd.Series(importances, index=feat_names).
    sort_values(ascending=True)

# Plot
feat_imp_df.plot(kind='barh', title='Feature Importances (Random Forest)',
                 color='green')
plt.xlabel('Importance Score')
plt.show()
```



```
[51]: from sklearn.tree import plot_tree

# Visualize one decision tree from the forest
plt.figure(figsize=(20, 10))
plot_tree(rf.estimators_[0], feature_names=feat_names, class_names=rf.classes_,
          filled=True)
plt.title("Decision Tree from Random Forest")
plt.show()
```



#### Q16.Boosting Methods

Implement AdaBoost and XGBoost on a classification task. Compare their accuracy and runtime performance.

```
[52]: import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import time

# Load Iris dataset
df = sns.load_dataset('iris')
X = df.drop('species', axis=1)
y = df['species']

# Encode target labels
y = y.astype('category').cat.codes # Converts species into numeric values

# Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪random_state=42)
```

```
[53]: from sklearn.ensemble import AdaBoostClassifier

# Train AdaBoost
start_time = time.time()
ada = AdaBoostClassifier(n_estimators=50, random_state=42)
ada.fit(X_train, y_train)
```

```

ada_time = time.time() - start_time

# Predict & evaluate
y_pred_ada = ada.predict(X_test)
acc_ada = accuracy_score(y_test, y_pred_ada)
print("AdaBoost Accuracy:", round(acc_ada * 100, 2), "%")
print("AdaBoost Training Time:", round(ada_time, 4), "seconds")

```

AdaBoost Accuracy: 93.33 %  
AdaBoost Training Time: 0.1158 seconds

```

[54]: from xgboost import XGBClassifier

# Train XGBoost
start_time = time.time()
xgb = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss')
xgb.fit(X_train, y_train)
xgb_time = time.time() - start_time

# Predict & evaluate
y_pred_xgb = xgb.predict(X_test)
acc_xgb = accuracy_score(y_test, y_pred_xgb)
print("XGBoost Accuracy:", round(acc_xgb * 100, 2), "%")
print("XGBoost Training Time:", round(xgb_time, 4), "seconds")

```

```

-----
ModuleNotFoundError                                Traceback (most recent call last)
Cell In[54], line 1
----> 1 from xgboost import XGBClassifier
      3 # Train XGBoost
      4 start_time = time.time()

ModuleNotFoundError: No module named 'xgboost'

```

!pip install xgboost

```

[57]: from xgboost import XGBClassifier

# Train XGBoost
start_time = time.time()
xgb = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss')
xgb.fit(X_train, y_train)
xgb_time = time.time() - start_time

# Predict & evaluate
y_pred_xgb = xgb.predict(X_test)
acc_xgb = accuracy_score(y_test, y_pred_xgb)

```

```
print("XGBoost Accuracy:", round(acc_xgb * 100, 2), "%")
print("XGBoost Training Time:", round(xgb_time, 4), "seconds")
```

XGBoost Accuracy: 100.0 %  
XGBoost Training Time: 0.0595 seconds

C:\Users\Shinde Ankita\AppData\Local\Programs\Python\Python313\Lib\site-packages\xgboost\training.py:183: UserWarning: [13:03:43] WARNING: C:\actions-runner\\_work\xgboost\xgboost\src\learner.cc:738:  
Parameters: { "use\_label\_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
```

Q17.K-Nearest Neighbors (KNN)

Implement a KNN classifier for a classification task. Experiment with different values of K and analyze their impact on accuracy.

```
[59]: import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score

# Load dataset
df = sns.load_dataset("iris")
X = df.drop("species", axis=1)
y = df["species"]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Feature scaling (important for KNN!)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
[60]: from sklearn.neighbors import KNeighborsClassifier
import matplotlib.pyplot as plt

k_values = list(range(1, 21))
accuracies = []

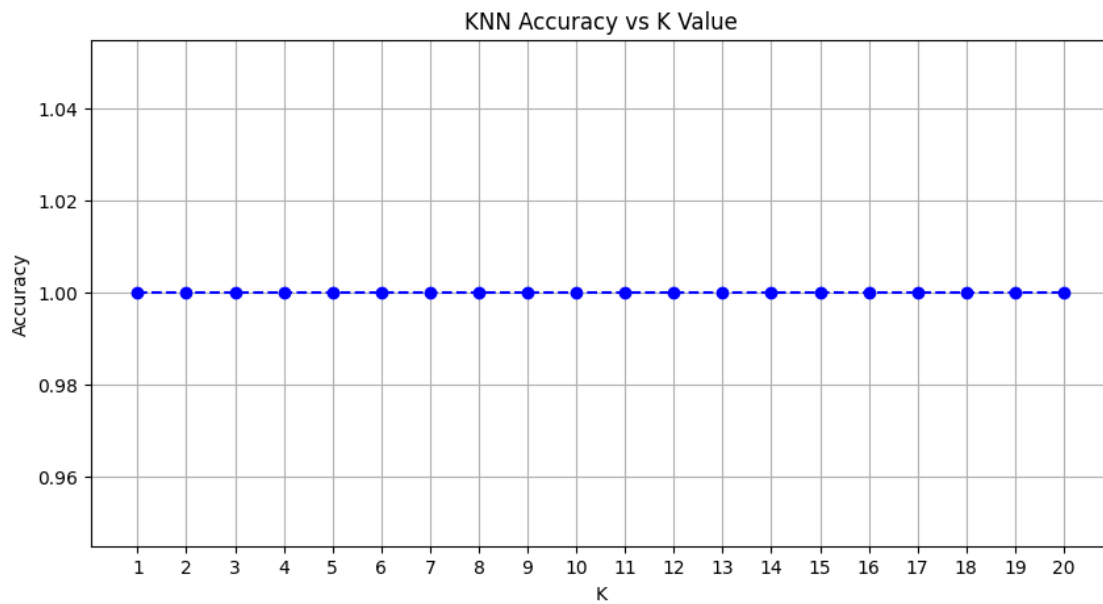
for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train_scaled, y_train)
    y_pred = knn.predict(X_test_scaled)
    acc = accuracy_score(y_test, y_pred)
```

```

        accuracies.append(acc)

# Plot accuracy vs K
plt.figure(figsize=(10, 5))
plt.plot(k_values, accuracies, marker='o', linestyle='--', color='blue')
plt.title('KNN Accuracy vs K Value')
plt.xlabel('K')
plt.ylabel('Accuracy')
plt.xticks(k_values)
plt.grid(True)
plt.show()

```



#### Q18.Support Vector Machines (SVM)

Train an SVM model with both linear and RBF kernels on a dataset. Visualize the decision boundaries.

```

[61]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score

# Load dataset
df = sns.load_dataset("iris")

```

```

# Use only 2 classes: setosa and versicolor
df = df[df['species'].isin(['setosa', 'versicolor'])]

# Use 2 features for visualization
X = df[['sepal_length', 'sepal_width']]
y = df['species']

# Encode labels
y = y.astype('category').cat.codes # 0=setosa, 1=versicolor

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)

# Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

```

```

[62]: # Linear kernel
svm_linear = SVC(kernel='linear')
svm_linear.fit(X_train_scaled, y_train)

# RBF kernel
svm_rbf = SVC(kernel='rbf')
svm_rbf.fit(X_train_scaled, y_train)

# Accuracy
acc_linear = accuracy_score(y_test, svm_linear.predict(X_test_scaled))
acc_rbf = accuracy_score(y_test, svm_rbf.predict(X_test_scaled))

print("Linear SVM Accuracy:", round(acc_linear * 100, 2), "%")
print("RBF SVM Accuracy:", round(acc_rbf * 100, 2), "%")

```

Linear SVM Accuracy: 100.0 %

RBF SVM Accuracy: 100.0 %

```

[63]: import numpy as np

def plot_decision_boundary(model, title):
    x_min, x_max = X_train_scaled[:, 0].min() - 1, X_train_scaled[:, 0].max() + 1
    ↪1
    y_min, y_max = X_train_scaled[:, 1].min() - 1, X_train_scaled[:, 1].max() + 1
    ↪1
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 200),
                          np.linspace(y_min, y_max, 200))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)


```

```


plt.figure(figsize=(8, 6))
plt.contourf(xx, yy, Z, alpha=0.3, cmap=plt.cm.coolwarm)
plt.scatter(X_train_scaled[:, 0], X_train_scaled[:, 1], c=y_train, cmap=plt.
cm.coolwarm, edgecolors='k')
plt.title(title)
plt.xlabel("Sepal Length (scaled)")
plt.ylabel("Sepal Width (scaled)")
plt.grid(True)
plt.show()

# Plot both
plot_decision_boundary(svm_linear, "SVM with Linear Kernel")
plot_decision_boundary(svm_rbf, "SVM with RBF Kernel")

```







### Q19.Regularization Techniques

Use Lasso and Ridge regression on a dataset. Analyze how they handle multicollinearity and reduce model complexity.

```
[65]: import numpy as np
import pandas as pd
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Lasso, Ridge
from sklearn.metrics import mean_squared_error

# Load the California housing dataset
housing = fetch_california_housing()
X = pd.DataFrame(housing.data, columns=housing.feature_names)
y = housing.target

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪random_state=42)
```

```

# Lasso Regression
lasso = Lasso(alpha=0.1) # Regularization parameter
lasso.fit(X_train, y_train)

# Ridge Regression
ridge = Ridge(alpha=0.1) # Regularization parameter
ridge.fit(X_train, y_train)

# Predictions
y_pred_lasso = lasso.predict(X_test)
y_pred_ridge = ridge.predict(X_test)

# Evaluate the models
lasso_rmse = np.sqrt(mean_squared_error(y_test, y_pred_lasso))
ridge_rmse = np.sqrt(mean_squared_error(y_test, y_pred_ridge))

print(f"Lasso RMSE: {lasso_rmse}")
print(f"Ridge RMSE: {ridge_rmse}")

# Compare the coefficients
print("\nLasso Coefficients:")
print(lasso.coef_)

print("\nRidge Coefficients:")
print(ridge.coef_)

```

Lasso RMSE: 0.7832697618354822

Ridge RMSE: 0.7455754517896762

Lasso Coefficients:

```
[ 3.92693362e-01  1.50810624e-02 -0.00000000e+00  0.00000000e+00
  1.64168387e-05 -3.14918929e-03 -1.14291203e-01 -9.93076483e-02]
```

Ridge Coefficients:

```
[ 4.48658477e-01  9.72442833e-03 -1.23292361e-01  7.82971747e-01
 -2.02924019e-06 -3.52627239e-03 -4.19791946e-01 -4.33705352e-01]
```

Q20.Introduction to Neural Networks

Build a simple Artificial Neural Network (ANN) to classify data. Use optimization algorithms (Gradient Descent, SGD) and visualize the loss during training.

```
[67]: !pip install torch torchvision matplotlib
```

Collecting torch

Downloading torch-2.6.0-cp313-cp313-win\_amd64.whl.metadata (28 kB)

Collecting torchvision

Downloading torchvision-0.21.0-cp313-cp313-win\_amd64.whl.metadata (6.3 kB)

Requirement already satisfied: matplotlib in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (3.10.1)

Collecting filelock (from torch)

Downloading filelock-3.18.0-py3-none-any.whl.metadata (2.9 kB)

Requirement already satisfied: typing-extensions>=4.10.0 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from torch) (4.13.2)

Collecting networkx (from torch)

Downloading networkx-3.4.2-py3-none-any.whl.metadata (6.3 kB)

Requirement already satisfied: Jinja2 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from torch) (3.1.6)

Collecting fsspec (from torch)

Downloading fsspec-2025.3.2-py3-none-any.whl.metadata (11 kB)

Requirement already satisfied: setuptools in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from torch) (78.1.0)

Collecting sympy==1.13.1 (from torch)

Downloading sympy-1.13.1-py3-none-any.whl.metadata (12 kB)

Collecting mpmath<1.4,>=1.1.0 (from sympy==1.13.1->torch)

Downloading mpmath-1.3.0-py3-none-any.whl.metadata (8.6 kB)

Requirement already satisfied: numpy in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from torchvision) (2.2.4)

Requirement already satisfied: pillow!=8.3.\*,>=5.3.0 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from torchvision) (11.1.0)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib) (1.3.1)

Requirement already satisfied: cycler>=0.10 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib) (4.57.0)

Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib) (1.4.8)

Requirement already satisfied: packaging>=20.0 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib) (24.2)

Requirement already satisfied: pyparsing>=2.3.1 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib) (3.2.3)

Requirement already satisfied: python-dateutil>=2.7 in c:\users\shinde ankita\appdata\local\programs\python\python313\lib\site-packages (from matplotlib) (2.9.0.post0)

Requirement already satisfied: six>=1.5 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from python-  
dateutil>=2.7->matplotlib) (1.17.0)

Requirement already satisfied: MarkupSafe>=2.0 in c:\users\shinde  
ankita\appdata\local\programs\python\python313\lib\site-packages (from  
jinja2->torch) (3.0.2)

Downloading torch-2.6.0-cp313-cp313-win\_amd64.whl (204.1 MB)

```
----- 0.0/204.1 MB ? eta -:-:--
----- 0.5/204.1 MB 4.9 MB/s eta 0:00:42
----- 2.1/204.1 MB 7.2 MB/s eta 0:00:29
----- 3.9/204.1 MB 7.6 MB/s eta 0:00:27
- ----- 5.2/204.1 MB 7.3 MB/s eta 0:00:28
- ----- 6.8/204.1 MB 7.5 MB/s eta 0:00:27
- ----- 8.7/204.1 MB 7.7 MB/s eta 0:00:26
-- ----- 10.2/204.1 MB 7.6 MB/s eta 0:00:26
-- ----- 11.8/204.1 MB 7.7 MB/s eta 0:00:25
-- ----- 13.6/204.1 MB 7.8 MB/s eta 0:00:25
-- ----- 15.2/204.1 MB 7.8 MB/s eta 0:00:25
--- ----- 17.0/204.1 MB 7.8 MB/s eta 0:00:24
--- ----- 18.6/204.1 MB 7.9 MB/s eta 0:00:24
---- ----- 20.4/204.1 MB 7.9 MB/s eta 0:00:24
---- ----- 22.3/204.1 MB 8.0 MB/s eta 0:00:23
---- ----- 24.1/204.1 MB 8.0 MB/s eta 0:00:23
---- ----- 25.4/204.1 MB 7.9 MB/s eta 0:00:23
---- ----- 27.3/204.1 MB 7.9 MB/s eta 0:00:23
----- ----- 28.8/204.1 MB 7.9 MB/s eta 0:00:23
----- ----- 30.7/204.1 MB 7.9 MB/s eta 0:00:22
----- ----- 32.5/204.1 MB 8.0 MB/s eta 0:00:22
----- ----- 34.1/204.1 MB 8.0 MB/s eta 0:00:22
----- ----- 35.7/204.1 MB 8.0 MB/s eta 0:00:22
----- ----- 37.0/204.1 MB 7.9 MB/s eta 0:00:22
----- ----- 38.8/204.1 MB 7.9 MB/s eta 0:00:21
----- ----- 40.4/204.1 MB 7.9 MB/s eta 0:00:21
----- ----- 41.9/204.1 MB 7.9 MB/s eta 0:00:21
----- ----- 43.8/204.1 MB 7.9 MB/s eta 0:00:21
----- ----- 45.4/204.1 MB 7.9 MB/s eta 0:00:21
----- ----- 47.2/204.1 MB 7.9 MB/s eta 0:00:20
----- ----- 48.5/204.1 MB 7.9 MB/s eta 0:00:20
----- ----- 50.3/204.1 MB 7.9 MB/s eta 0:00:20
----- ----- 51.4/204.1 MB 7.9 MB/s eta 0:00:20
----- ----- 52.2/204.1 MB 7.7 MB/s eta 0:00:20
----- ----- 53.5/204.1 MB 7.7 MB/s eta 0:00:20
----- ----- 55.1/204.1 MB 7.6 MB/s eta 0:00:20
----- ----- 56.6/204.1 MB 7.6 MB/s eta 0:00:20
----- ----- 58.2/204.1 MB 7.6 MB/s eta 0:00:20
----- ----- 60.0/204.1 MB 7.6 MB/s eta 0:00:19
----- ----- 61.9/204.1 MB 7.6 MB/s eta 0:00:19
----- ----- 63.4/204.1 MB 7.7 MB/s eta 0:00:19
```

```

----- 65.0/204.1 MB 7.6 MB/s eta 0:00:19
----- 66.8/204.1 MB 7.7 MB/s eta 0:00:18
----- 68.4/204.1 MB 7.7 MB/s eta 0:00:18
----- 70.0/204.1 MB 7.7 MB/s eta 0:00:18
----- 71.8/204.1 MB 7.7 MB/s eta 0:00:18
----- 73.7/204.1 MB 7.7 MB/s eta 0:00:17
----- 75.5/204.1 MB 7.7 MB/s eta 0:00:17
----- 76.8/204.1 MB 7.7 MB/s eta 0:00:17
----- 77.1/204.1 MB 7.7 MB/s eta 0:00:17
----- 78.1/204.1 MB 7.5 MB/s eta 0:00:17
----- 79.7/204.1 MB 7.5 MB/s eta 0:00:17
----- 81.5/204.1 MB 7.6 MB/s eta 0:00:17
----- 83.1/204.1 MB 7.6 MB/s eta 0:00:16
----- 84.7/204.1 MB 7.6 MB/s eta 0:00:16
----- 86.2/204.1 MB 7.6 MB/s eta 0:00:16
----- 87.8/204.1 MB 7.6 MB/s eta 0:00:16
----- 89.7/204.1 MB 7.6 MB/s eta 0:00:16
----- 91.5/204.1 MB 7.6 MB/s eta 0:00:15
----- 93.1/204.1 MB 7.6 MB/s eta 0:00:15
----- 94.9/204.1 MB 7.6 MB/s eta 0:00:15
----- 96.5/204.1 MB 7.6 MB/s eta 0:00:15
----- 98.3/204.1 MB 7.6 MB/s eta 0:00:14
----- 99.9/204.1 MB 7.6 MB/s eta 0:00:14
----- 101.7/204.1 MB 7.7 MB/s eta 0:00:14
----- 103.3/204.1 MB 7.7 MB/s eta 0:00:14
----- 104.9/204.1 MB 7.7 MB/s eta 0:00:13
----- 106.7/204.1 MB 7.7 MB/s eta 0:00:13
----- 107.0/204.1 MB 7.7 MB/s eta 0:00:13
----- 108.5/204.1 MB 7.6 MB/s eta 0:00:13
----- 110.4/204.1 MB 7.6 MB/s eta 0:00:13
----- 112.2/204.1 MB 7.6 MB/s eta 0:00:13
----- 114.0/204.1 MB 7.6 MB/s eta 0:00:12
----- 115.6/204.1 MB 7.6 MB/s eta 0:00:12
----- 117.4/204.1 MB 7.6 MB/s eta 0:00:12
----- 119.0/204.1 MB 7.6 MB/s eta 0:00:12
----- 120.8/204.1 MB 7.6 MB/s eta 0:00:11
----- 122.7/204.1 MB 7.7 MB/s eta 0:00:11
----- 124.3/204.1 MB 7.7 MB/s eta 0:00:11
----- 125.8/204.1 MB 7.7 MB/s eta 0:00:11
----- 127.4/204.1 MB 7.7 MB/s eta 0:00:11
----- 129.2/204.1 MB 7.7 MB/s eta 0:00:10
----- 130.8/204.1 MB 7.7 MB/s eta 0:00:10
----- 132.6/204.1 MB 7.7 MB/s eta 0:00:10
----- 134.2/204.1 MB 7.7 MB/s eta 0:00:10
----- 135.8/204.1 MB 7.7 MB/s eta 0:00:09
----- 137.6/204.1 MB 7.7 MB/s eta 0:00:09
----- 139.5/204.1 MB 7.7 MB/s eta 0:00:09
----- 141.0/204.1 MB 7.7 MB/s eta 0:00:09

```

```

----- 142.9/204.1 MB 7.7 MB/s eta 0:00:08
----- 144.4/204.1 MB 7.7 MB/s eta 0:00:08
----- 146.3/204.1 MB 7.7 MB/s eta 0:00:08
----- 147.6/204.1 MB 7.7 MB/s eta 0:00:08
----- 148.9/204.1 MB 7.7 MB/s eta 0:00:08
----- 150.2/204.1 MB 7.7 MB/s eta 0:00:08
----- 151.8/204.1 MB 7.7 MB/s eta 0:00:07
----- 153.4/204.1 MB 7.7 MB/s eta 0:00:07
----- 154.7/204.1 MB 7.7 MB/s eta 0:00:07
----- 156.0/204.1 MB 7.6 MB/s eta 0:00:07
----- 157.3/204.1 MB 7.6 MB/s eta 0:00:07
----- 157.3/204.1 MB 7.6 MB/s eta 0:00:07
----- 158.3/204.1 MB 7.5 MB/s eta 0:00:07
----- 159.4/204.1 MB 7.5 MB/s eta 0:00:06
----- 161.2/204.1 MB 7.5 MB/s eta 0:00:06
----- 162.8/204.1 MB 7.5 MB/s eta 0:00:06
----- 164.4/204.1 MB 7.5 MB/s eta 0:00:06
----- 166.2/204.1 MB 7.5 MB/s eta 0:00:06
----- 167.8/204.1 MB 7.5 MB/s eta 0:00:05
----- 169.6/204.1 MB 7.5 MB/s eta 0:00:05
----- 171.2/204.1 MB 7.5 MB/s eta 0:00:05
----- 172.8/204.1 MB 7.5 MB/s eta 0:00:05
----- 174.6/204.1 MB 7.5 MB/s eta 0:00:04
----- 176.2/204.1 MB 7.5 MB/s eta 0:00:04
----- 178.0/204.1 MB 7.5 MB/s eta 0:00:04
----- 179.8/204.1 MB 7.5 MB/s eta 0:00:04
----- 181.7/204.1 MB 7.6 MB/s eta 0:00:03
----- 183.0/204.1 MB 7.5 MB/s eta 0:00:03
----- 184.8/204.1 MB 7.6 MB/s eta 0:00:03
----- 186.4/204.1 MB 7.6 MB/s eta 0:00:03
----- 188.2/204.1 MB 7.6 MB/s eta 0:00:03
----- 190.1/204.1 MB 7.6 MB/s eta 0:00:02
----- 191.6/204.1 MB 7.6 MB/s eta 0:00:02
----- 192.9/204.1 MB 7.6 MB/s eta 0:00:02
----- 194.8/204.1 MB 7.6 MB/s eta 0:00:02
----- 196.3/204.1 MB 7.6 MB/s eta 0:00:02
----- 198.2/204.1 MB 7.6 MB/s eta 0:00:01
----- 200.0/204.1 MB 7.6 MB/s eta 0:00:01
----- 201.9/204.1 MB 7.6 MB/s eta 0:00:01
----- 203.2/204.1 MB 7.6 MB/s eta 0:00:01
----- 203.9/204.1 MB 7.6 MB/s eta 0:00:01
----- 204.1/204.1 MB 7.5 MB/s eta 0:00:00
Downloading sympy-1.13.1-py3-none-any.whl (6.2 MB)
----- 0.0/6.2 MB ? eta -:--:--
----- 1.6/6.2 MB 8.5 MB/s eta 0:00:01
----- 3.4/6.2 MB 8.5 MB/s eta 0:00:01
----- 5.2/6.2 MB 8.4 MB/s eta 0:00:01
----- 6.2/6.2 MB 7.8 MB/s eta 0:00:00

```

```

Downloading torchvision-0.21.0-cp313-cp313-win_amd64.whl (1.6 MB)
----- 0.0/1.6 MB ? eta -:--:--
----- 1.3/1.6 MB 8.2 MB/s eta 0:00:01
----- 1.6/1.6 MB 7.0 MB/s eta 0:00:00
Downloading filelock-3.18.0-py3-none-any.whl (16 kB)
Downloading fsspec-2025.3.2-py3-none-any.whl (194 kB)
Downloading networkx-3.4.2-py3-none-any.whl (1.7 MB)
----- 0.0/1.7 MB ? eta -:--:--
----- 1.6/1.7 MB 8.7 MB/s eta 0:00:01
----- 1.7/1.7 MB 8.3 MB/s eta 0:00:00
Downloading mpmath-1.3.0-py3-none-any.whl (536 kB)
----- 0.0/536.2 kB ? eta -:--:--
----- 536.2/536.2 kB 7.6 MB/s eta 0:00:00
Installing collected packages: mpmath, sympy, networkx, fsspec, filelock, torch,
torchvision
Successfully installed filelock-3.18.0 fsspec-2025.3.2 mpmath-1.3.0
networkx-3.4.2 sympy-1.13.1 torch-2.6.0 torchvision-0.21.0

```

```

[75]: import torch
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
import numpy as np

```

```

[76]: # Load the Iris dataset
data = load_iris()
X = data.data
y = data.target

# Standardize the dataset
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

# Convert to PyTorch tensors
X_train = torch.tensor(X_train, dtype=torch.float32)
X_test = torch.tensor(X_test, dtype=torch.float32)
y_train = torch.tensor(y_train, dtype=torch.long)
y_test = torch.tensor(y_test, dtype=torch.long)

```

```
[77]: class SimpleANN(nn.Module):
    def __init__(self):
        super(SimpleANN, self).__init__()
        # Define the layers
        self.layer1 = nn.Linear(4, 10) # Input layer (4 features to 10 neurons
    ↪in the hidden layer)
        self.layer2 = nn.Linear(10, 3) # Output layer (10 neurons to 3 output
    ↪classes)

    def forward(self, x):
        x = torch.relu(self.layer1(x)) # ReLU activation function for hidden
    ↪layer
        x = self.layer2(x) # Output layer (logits)
        return x
```

```
[78]: # Initialize the model
model = SimpleANN()

# Loss function (Cross Entropy for multi-class classification)
criterion = nn.CrossEntropyLoss()

# Optimizer: Gradient Descent (SGD can be used as well)
optimizer = optim.SGD(model.parameters(), lr=0.01)

# For visualization, we will store the loss values
train_losses = []
```

```
[79]: epochs = 200
for epoch in range(epochs):
    model.train() # Set the model to training mode
    optimizer.zero_grad() # Zero the gradients from previous step

    # Forward pass
    outputs = model(X_train)
    loss = criterion(outputs, y_train) # Calculate the loss

    # Backward pass and optimization
    loss.backward()
    optimizer.step()

    # Store the loss for visualization
    train_losses.append(loss.item())

    if (epoch + 1) % 20 == 0:
        print(f'Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}')
```

Epoch [20/200], Loss: 1.0315




```
Epoch [40/200], Loss: 0.9636
Epoch [60/200], Loss: 0.9111
Epoch [80/200], Loss: 0.8685
Epoch [100/200], Loss: 0.8327
Epoch [120/200], Loss: 0.8017
Epoch [140/200], Loss: 0.7744
Epoch [160/200], Loss: 0.7499
Epoch [180/200], Loss: 0.7278
Epoch [200/200], Loss: 0.7075
```

```
[80]: # Set the model to evaluation mode
      model.eval()

      # Get the predictions on the test set
      with torch.no_grad():
          outputs = model(X_test)
          _, predicted = torch.max(outputs, 1) # Get the class with the highest
          ↪probability
          accuracy = accuracy_score(y_test, predicted)
          print(f'Accuracy on Test Set: {accuracy * 100:.2f}%')
```

Accuracy on Test Set: 73.33%

```
[81]: # Plot the training loss
      plt.plot(train_losses)
      plt.xlabel('Epochs')
      plt.ylabel('Loss')
      plt.title('Training Loss During Epochs')
      plt.show()
```



[ ]:

[ ]: